

APM Agents Node.JS vs Java

Compare Java and Node.JS tooling related to APM Agents



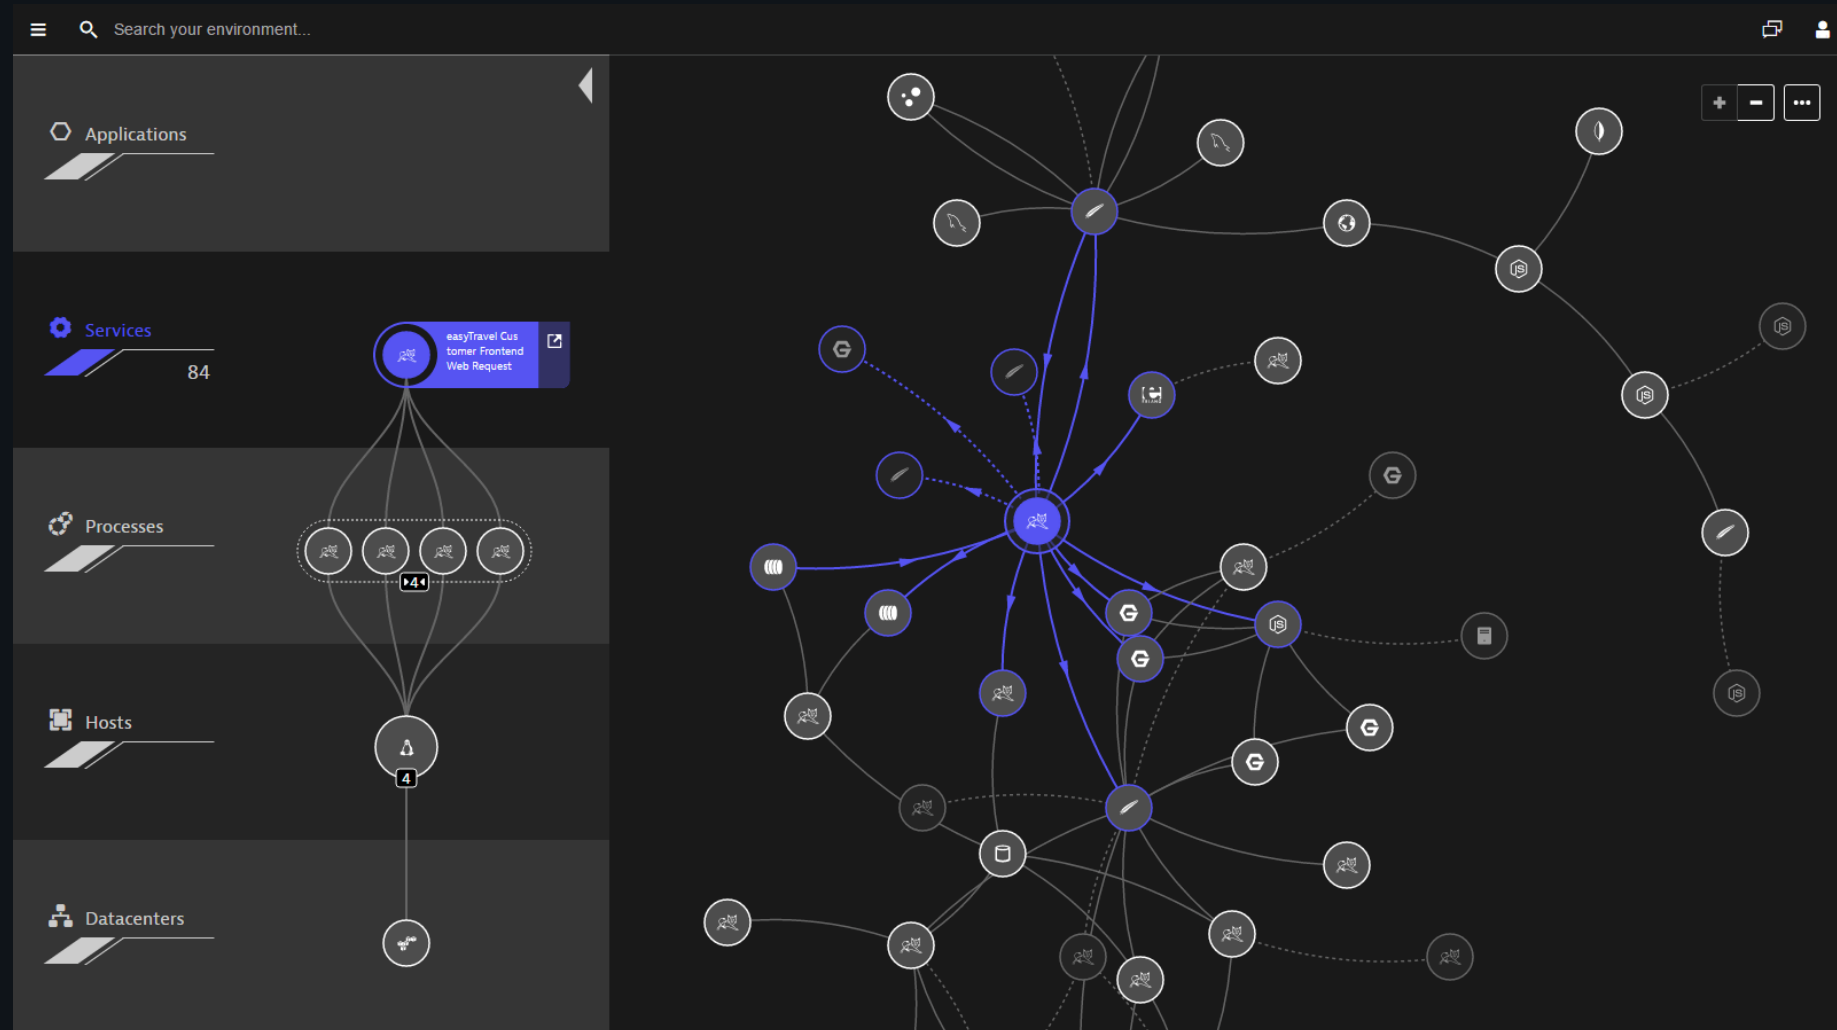


Introduction – APM

- Agents are injected into customer applications in production
- Agents continuously report data
 - Metrics like CPU/Memory usage, GC time, Event Loop timing,...
 - Transactions (focus on topology, duration and CPU time)
 - References to source code (stack traces, linked with transactions / background)
 - Low resolution CPU Profiling (hotspots on transactions / background)
- Agents add/extract tags to link transactions between processes ("Trace Context")
- Agent shall not change application behavior
- Servers correlate agent data and visualize it (e.g. topology, detected problems,...)



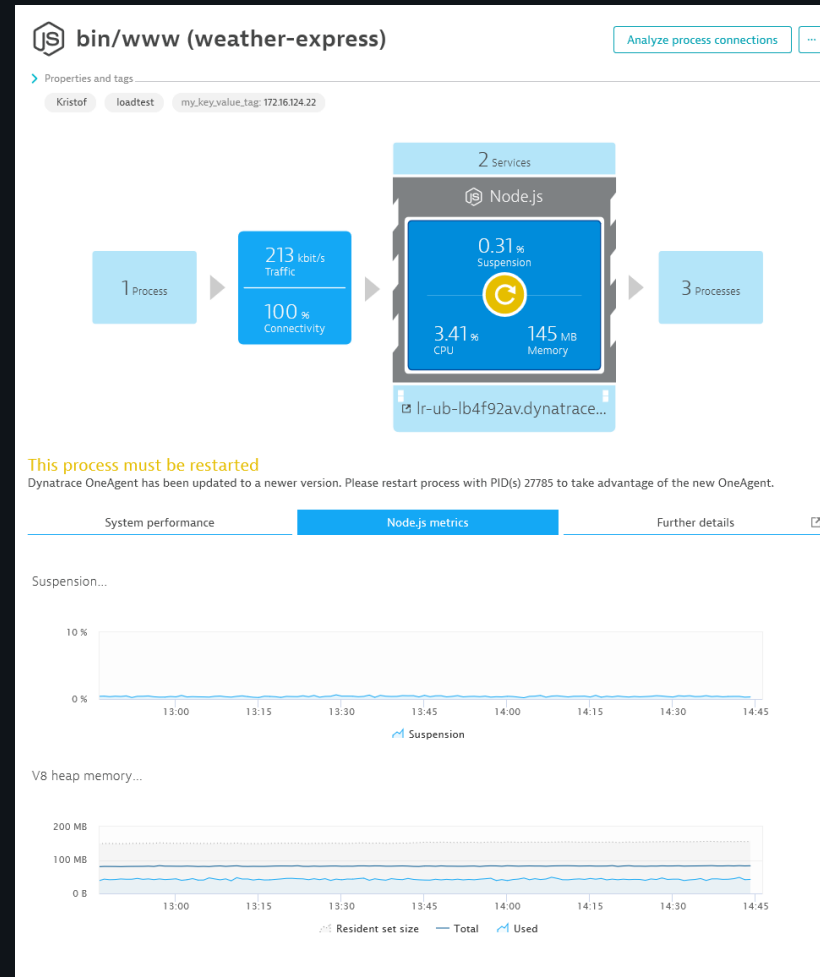
Infrastructure / Topology








Process Details including Metrics



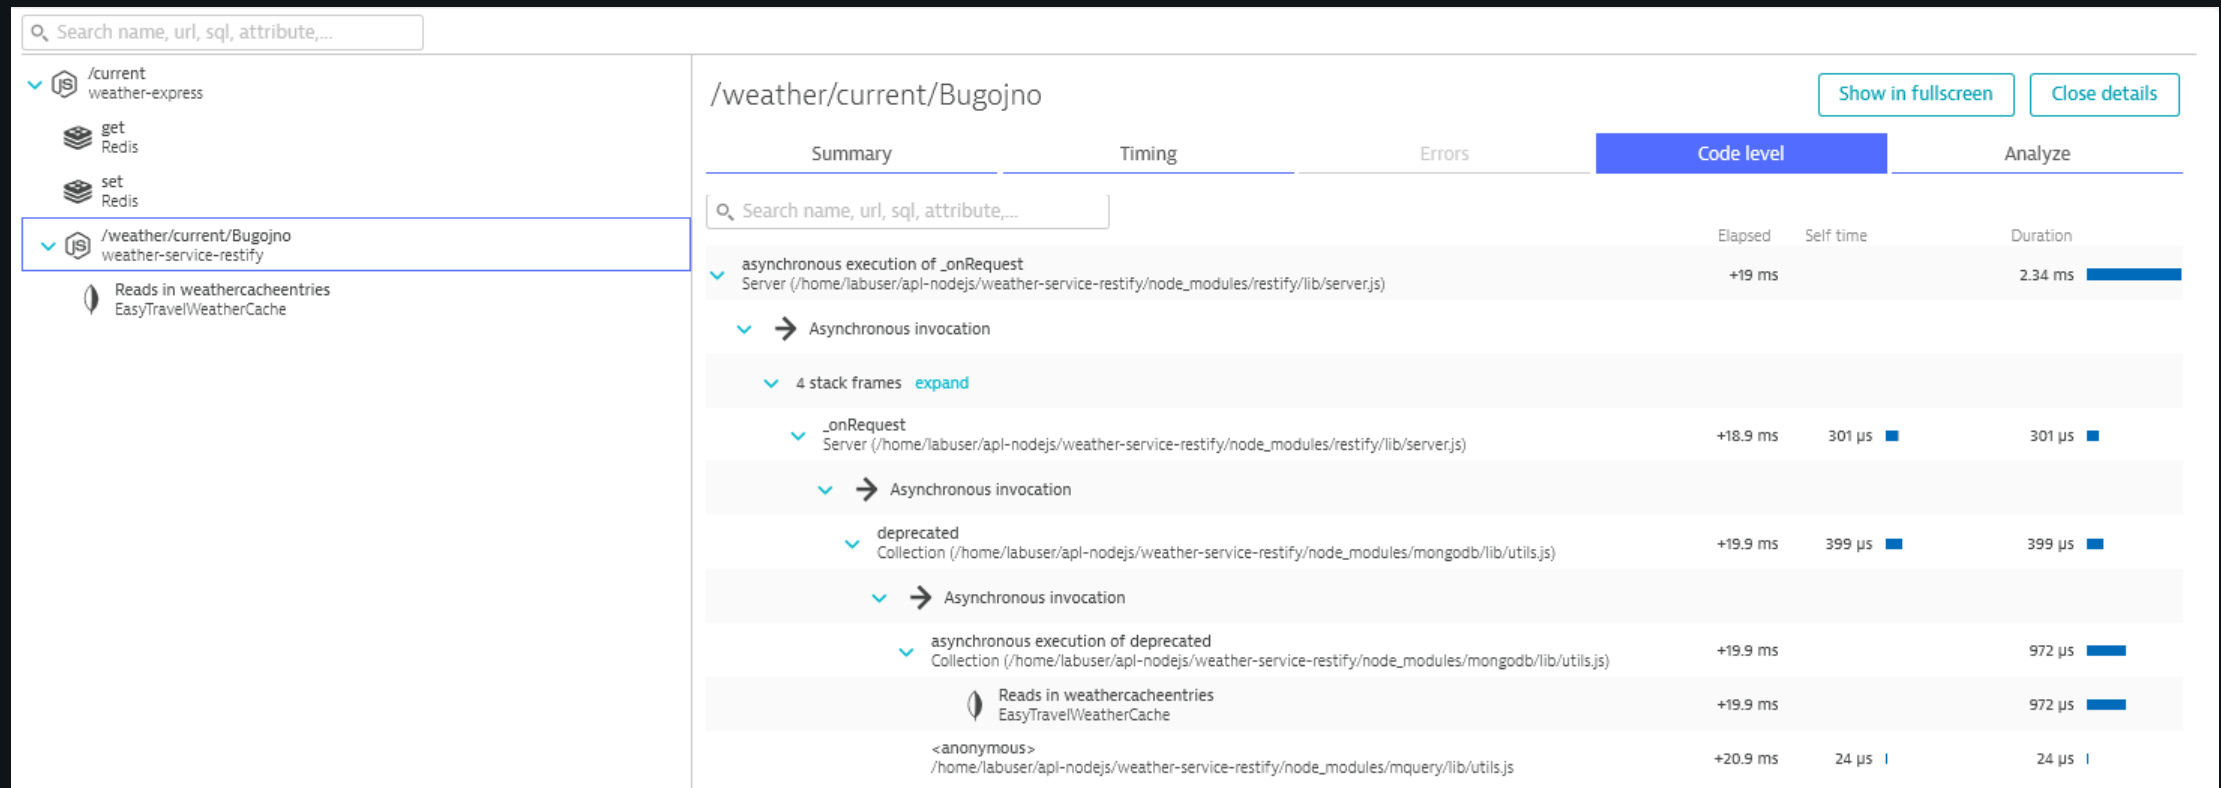


CPU Hotspots (Service / Background)





Transactions with code level visibility





Agent initialization/instrumentation

Node.JS	Java
No dedicated APIs to load an agent and instrument code	Standardized tooling API (JVMTI)
Loaded via <code>require()</code> as <i>first</i> statement in application (or via <code>-r</code> command line option)	Dedicated option (<code>-agentpath</code>) to load and initialize agent
Monkey Patch exports by intercepting <code>module._load()</code>	Modify Byte Code during class loading (JVMTI hook)
Internals/implementation details needed	Internals/implementation details needed
No generic database/metric interfaces	Generic interfaces for databases (JDBC) and metrics (JMX)
Use of public APIs to get e.g. metrics (GC, ...)	Use of public APIs to get e.g. metrics (GC, ...)
Native ABI stable within a major Node.JS version No diagnostics/debugging/... APIs in N-API	ABI stable since Java 5 (but extended since then)



Monkey Patching vs Byte Code modifications

Node.JS	Java
Monkey Patching is frequently used - also by a lot user space modules	Byte Code instrumentation is usually not used by application programmers
No access to internal core modules No access to file local functions Non configurable properties avoid monkey patching Private symbols, private fields, ... Exported functions may be cached inside the file ("partial" instrumentation)	Full access to all loaded classes including inner classes, private members,... Even hot modification is possible
Agent and application frames intermixed on call stack	Agent frames on call stack only as leaf



Code Level visibility / hotspot detection / Profiling

Node.JS	Java
No API to get a single JS call stack from a different thread	API available to capture a stack trace from any thread Allows to add context to stack traces (e.g. transaction) Allows to monitor overhead cause by agent
CPU Profiler included (start/stop, result is a tree)	Profiler can be implemented based on above API
Anonymous functions are very common but have mostly no context when seen on a call stack Inferred name is hard to get and sometimes misleading	Anonymous functions (lambdas) are always part of a class and usually implement an interface
Async IO hides a lot of the transactional context (e.g. no info about Cpu cycles used inside LibUv) Use of <code>nextTick/...</code> spreads CPU time/wait time Async Hooks available for tracking (at least in core)	IO often synchronous therefore stack sampling shows where time is spent no framework similar to Async Hooks (instrumentation of threading APIs needed – part of JVM)



Memory Profiling, Stack Traces

Node.JS	Java
Complete Memory Profiler included (Sampling Heap Profiler intended for production use)	JVMTI API to install a sync callback for allocations after every ~x Bytes (since Java 11)
No API to iterate function objects on stacks (strict mode) Function identities are lost as name/script/line can't be correlated to the concrete function instance	Iterate function/classes objects on stacks



Summary

- Java offers standardized and more generic low level APIs with focus on tool development
- Node.JS offers ready to use functionality for specific use cases
- Strongly typed nature of Java offers benefits to identify/map functions in source code
- Sync nature of a lot Java APIs results in easier monitoring/correlation
- Production readiness (Stability/Overhead) not clearly documented for some Node.JS features
- Node.JS core is much smaller than Java, a lot userland code with duplications (e.g. Promise libs) and independent versioning

Questions ???

Software intelligence for the enterprise cloud





Software intelligence built for the enterprise cloud

Go beyond APM with the Dynatrace all-in-one platform





Better data makes Dynatrace A.I. and massive automation possible

High fidelity data

Mapped end-to-end

Deterministic AI

Answers + Action



Automated problem detection
Business impact determined
Root cause explained
No alert storms
Trigger self healing

Completely automated

Dynatrace - Software intelligence built for the enterprise cloud

Application performance management

Microservices & containers
Lifecycle analytics

Code-level
Transaction tracing

Database monitoring
Monolith and mainframe

Cloud infrastructure monitoring

Cloud monitoring
Container monitoring

Virtualization monitoring
Server monitoring

Network monitoring
Log monitoring

AIOps

Continuous auto-discovery
Automatic topology

Anomaly detection
Root cause analysis

Prediction
Third party data & metrics

Digital experience management

Real user monitoring
Session replay

Mobile app monitoring
Synthetic monitoring

RUM for SaaS vendors
Digital experience insights

All-in-one software intelligence platform.

Think you know Dynatrace? Think again.



550+ engineers in R&D make it all possible

66% more than nearest
competitor

New releases every 2 weeks

26 releases per year



A trusted leader

8 year leader in Gartner MQ

#1 market share leader

Trusted by 72 out of the
Fortune 100

Get ready to be
amazed in 5
minutes or less

dynatrace.com/trial



OpenShift
Primed Partner

Pivotal

Technology
partner



Docker Certified



AWS Advanced
Technology Partner



Azure Partner



dynatrace.com